

CASE STUDY 1 CUSTOMER TRANSACTION

In this Case study, we have two tables **customer** and **transactions** in database **acadgilddb** as shown below in screenshot:

```
acadgild@localhost:~  
hive> use acadgilddb ;  
OK  
Time taken: 0.025 seconds  
hive> show tables;  
OK  
customer  
transactions  
Time taken: 0.042 seconds, Fetched: 2 row(s)  
hive>
```

customer table have five columns consist of **customer ID**, **customer first name**, **customer last name**, **age** and **customer profession**.

We can find customer schema by typing: **describe customer** as shown below:

```
acadgild@localhost:~  
hive> describe customer;  
OK  
custid          int  
fname           string  
lname           string  
age             int  
profession       string  
Time taken: 0.07 seconds, Fetched: 5 row(s)  
hive>
```

Below screen shot shows the data present in the **customer** table

```
acadgild@localhost:~  
hive> select * from customer;  
OK  
101    Amitabh Bacchan 65      Actor  
102    Sharukh Khan   45      Doctor  
103    Akshay Kumar   38      Dentist  
104    Anubahv kumar  58      Business  
105    Pawan Trivedi  34      service  
106    Aamir Null     42      scientest  
107    Salman Khan    43      Surgen  
108    Ranbir Kapoor  26      Industrialist  
Time taken: 0.207 seconds, Fetched: 8 row(s)  
hive>
```

transaction table have nine columns consist of transaction **number**, **transaction date**, **customer ID**, **amount**,**category**,**product detail**,**city**,**state**,**spendby details**.

We will find this detail about table by: “**describe transaction**” as shown below.

```
acadgild@localhost:~
hive> describe transactions;
OK
txnno          int
txndate        string
custno         int
amount         double
category       string
product        string
city           string
state          string
spendby        string
Time taken: 0.068 seconds, Fetched: 9 row(s)
hive>
```

Below screen shot shows the **transactions** table data.

```
acadgild@localhost:~
hive> select * from transactions;
OK
97834 05/02/2018 101 965.0 Entertainment Movie Pune Maharashtra Daughter
98396 12/01/2018 102 239.0 Food Grocery Patna Bihar Self
34908 06/01/2018 101 875.0 Travel Air Bangalore Karnataka Spouse
70958 17/02/2018 104 439.0 Food Restaurant Delhi Delhi Wife
9874 21/01/2018 105 509.0 Entertainment Park Kolkata West Bengal NULL
94585 19/01/2018 106 629.0 Rent House Hyderabad Telangana Self
45509 20/01/2018 107 953.0 Travel Rail Chennai Tamil Nadu Brother
7864 01/02/2018 108 569.0 Rent Parking Goa Goa Wife
Time taken: 0.165 seconds, Fetched: 8 row(s)
hive>
```

Let us solve the following use cases using these tables:-

1. Find out the number of transaction done by each customer.

To find the number of transaction done by each customer we have to use following query to get the result:-

SELECT custno, **COUNT**(*) as totaltrx **FROM** transactions

GROUP BY custno;

```

acadmild@localhost:~
hive> select custno,count(*) as totaltrx from TRANSACTIONS
> group by custno;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a diff
ngine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180424172440_8095a9c0-cef5-409b-b10a-bdd7e22043ad
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1524549565701_0019, Tracking URL = http://localhost:8088/proxy/application_1524549565701_0019/
Kill Command = /home/acadmild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1524549565701_0019
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 1
2018-04-24 17:24:47,994 Stage-1 map = 0%, reduce = 0%
2018-04-24 17:24:54,341 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.96 sec
2018-04-24 17:25:01,767 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 4.7 sec
MapReduce Total cumulative CPU time: 4 seconds 700 msec
Ended Job = job_1524549565701_0019
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 4.7 sec HDFS Read: 9719 HDFS Write: 213 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 700 msec
OK
101      2
102      1
104      1
105      1
106      1
107      1
108      1
Time taken: 22.082 seconds, Fetched: 7 row(s)
hive>

```

Above screen shot shows the output with number of transaction and customer ID.

We can also find the number of transaction done by each customer by getting the name of the customer by using join query as shown below

```

SELECT t1.id, t1.f, t1.l, COUNT (t1.txn) FROM

(SELECT c.custid as id, c.fname as f, c.lname as l, t.txnno as txn

FROM customer c JOIN transactions t ON c.custid = t.custno) t1

GROUP BY t1.id, t1.f, t1.l;

```

```
acadmild@localhost:~
hive> SELECT t1.id, t1.f, t1.l, COUNT(t1.txn)
> FROM
> (SELECT c.custid as id, c.fname as f, c.lname as l, t.txnno as txn
> FROM customer c JOIN transactions t ON c.custid = t.custno) t1
> GROUP BY t1.id, t1.f, t1.l;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadmild_20180424163610_939ed1b9-a6d0-49a1-ab9a-346e230f58c7
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadmild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadmild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-04-24 16:36:18 Starting to launch local task to process map join; maximum memory = 477626368
2018-04-24 16:36:20 Dump the side-table for tag: 0 with group count: 8 into file: file:/tmp/acadmild/3b33eea5-b37a-4521-8229-b94f684d9dbf/hive_2018-04-24_16-36-10_910_8190550956199868160-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile40--.hashtable
2018-04-24 16:36:20 Uploaded 1 File to: file:/tmp/acadmild/3b33eea5-b37a-4521-8229-b94f684d9dbf/hive_2018-04-24_16-36-10_910_8190550956199868160-1/-local-10005/HashTable-Stage-2/MapJoin-mapfile40--.hashtable (522 bytes)
2018-04-24 16:36:20 End of local task; Time Taken: 2.088 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1524549565701_0016, Tracking URL = http://localhost:8088/proxy/application_1524549565701_0016/

Starting Job = job_1524549565701_0016, Tracking URL = http://localhost:8088/proxy/application_1524549565701_0016/
Kill Command = /home/acadmild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1524549565701_0016
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-04-24 16:36:28,961 Stage-2 map = 0%, reduce = 0%
2018-04-24 16:36:35,499 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 3.19 sec
2018-04-24 16:36:42,930 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 5.9 sec
MapReduce Total cumulative CPU time: 5 seconds 900 msec
Ended Job = job_1524549565701_0016
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 5.9 sec HDFS Read: 13925 HDFS Write: 307 SUCCESS
Total MapReduce CPU Time Spent: 5 seconds 900 msec
OK
101 Amitabh Bacchan 2
102 Sharukh Khan 1
104 Anubhav kumar 1
105 Pawan Trivedi 1
106 Aamir Null 1
107 Salman Khan 1
108 Ranbir Kapoor 1
Time taken: 33.139 seconds, Fetched: 7 row(s)
```

In above screen-shot we are able to see the output containing transaction done by each customer with their **first name and last name and customer ID.**

2. Create a new table called TRANSACTIONS_COUNT. This table should have three fields - custid, fname and count.

To solve above problem we have to use below query to create the table

CREATE TABLE transactions_Count

(custid INT, fname STRING, txn_count INT)

ROW FORMAT DELIMITED FIELDS TERMINATED by ',';

```

acadgild@localhost:~
hive> CREATE TABLE transactions_Count
> ( custid INT,
>   fname STRING,
>   txn_count INT
> )
> ROW FORMAT DELIMITED FIELDS TERMINATED BY ',';
OK
Time taken: 0.252 seconds
hive>

```

3. Now write a hive query in such a way that the query populates the data obtained in Step 1 above and populate the table in step 2 above.

To solve above problem we have to use **insert** query to insert data obtained from the problem number 2 into Transactions_count

INSERT INTO transactions_Count

SELECT t1.id, t1.f, COUNT(t1.txn) **FROM**

(**SELECT** c.custid as id, c.fname as f, t.txnno as txn

FROM customer c **JOIN** transactions t **ON** c.custid = t.custno)t1

GROUP BY t1.id, t1.f;

```

acadgild@localhost:~
hive> INSERT INTO transactions_Count
> SELECT t1.id, t1.f, COUNT(t1.txn)
> FROM
> (SELECT c.custid as id, c.fname as f, t.txnno as txn
>   FROM customer c JOIN transactions t ON c.custid = t.custno)t1
> GROUP BY t1.id, t1.f;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20180424170008_0582be7e-21b9-49d3-a9c4-9778a0033217
Total jobs = 1
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/acadgild/install/hive/apache-hive-2.3.2-bin/lib/log4j-slf4j-impl-2.6.2.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/acadgild/install/hadoop/hadoop-2.6.5/share/hadoop/common/lib/slf4j-log4j12-1.7.5.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
2018-04-24 17:00:15 Starting to launch local task to process map join: maximum memory = 477626368
2018-04-24 17:00:17 Dump the side-table for tag: 0 with group count: 8 into file: file:/tmp/acadgild/3b33eea5-b37a-4521-8229-b94f684d9dbf/hive_2018-04-24_17-00-08_209_4610626828641175211-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile50--.hashtable
2018-04-24 17:00:17 Uploaded 1 File to: file:/tmp/acadgild/3b33eea5-b37a-4521-8229-b94f684d9dbf/hive_2018-04-24_17-00-08_209_4610626828641175211-1/-local-10003/HashTable-Stage-2/MapJoin-mapfile50--.hashtable (469 bytes)
2018-04-24 17:00:17 End of local task; Time Taken: 2.072 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks not specified. Estimated from input data size: 1
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>

```

```

In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1524549565701_0017, Tracking URL = http://localhost:8088/proxy/application_1524549565701_0017/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1524549565701_0017
Hadoop job information for Stage-2: number of mappers: 1; number of reducers: 1
2018-04-24 17:00:25,201 Stage-2 map = 0%, reduce = 0%
2018-04-24 17:00:32,702 Stage-2 map = 100%, reduce = 0%, Cumulative CPU 2.7 sec
2018-04-24 17:00:40,171 Stage-2 map = 100%, reduce = 100%, Cumulative CPU 6.8 sec
MapReduce Total cumulative CPU time: 6 seconds 800 msec
Ended Job = job_1524549565701_0017
Loading data to table acadgild.db.transactions_count
MapReduce Jobs Launched:
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 6.8 sec HDFS Read: 14126 HDFS Write: 177 SUCCESS
Total MapReduce CPU Time Spent: 6 seconds 800 msec
OK
Time taken: 34.65 seconds
hive> select * from transactions_Count;
OK
101    Amitabh 2
102    Sharukh 1
104    Anubhav 1
105    Pawan   1
106    Aamir   1
107    Salman  1
108    Ranbir  1
Time taken: 0.162 seconds, Fetched: 7 row(s)
hive>

```

Above screen, shot shows that data obtained from query in case1 has successfully inserted in table transactions_count.

Query “SELECT * **FROM** transactions_Count” shows the result.

4. Now lets make the TRANSACTIONS_COUNT table Hbase compliant. In the sense, use SerDes and Storage handler features of hive to change the TRANSACTIONS_COUNT table to be able to create a TRANSACTIONS table in Hbase.

CREATE TABLE TRANSACTIONS_HBase(userID STRING,username STRING, count_txn STRING)

STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'

WITH SERDEPROPERTIES ('hbase.columns.mapping' = ': key,stats:username,stats:count_txn')

TBLPROPERTIES ('hbase.table.name' = 'TRANSACTIONS');

```

acadgild@localhost:~
hive> CREATE TABLE TRANSACTIONS_HBase(userID STRING,username STRING, count_txn INT)
> STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
> WITH SERDEPROPERTIES ('hbase.columns.mapping' = ':key,stats:username,stats:count_txn')
> TBLPROPERTIES ('hbase.table.name' = 'TRANSACTIONS');
OK
Time taken: 1.438 seconds
hive>

```

```

acadgild@localhost:~
hbase(main):023:0> list
TABLE
TRANSACTIONS
bulktable
clicks
emp
emp420
employee
htest
test
8 row(s) in 0.0130 seconds

```

```

acadgild@localhost:~
hbase(main):024:0> scan 'TRANSACTIONS'
ROW                                COLUMN+CELL
0 row(s) in 0.0280 seconds

hbase(main):025:0> █

```

5. Now insert the data in TRANSACTIONS_Hbase table using the query in step-3 again, this should populate the Hbase TRANSACTIONS table automatically.

To solve above problem we use insert query to transfer data from TRANSACTIONS_COUNT into TRANSACTIONS_HBASE.

INSERT INTO TRANSACTIONS_HBase

SELECT * FROM TRANSACTIONS_COUNT;

Below screen shot shows the TRANSACTIONS COUNT data is successfully inserted into TRANSACTIONS_HBase

```

acadgild@localhost:~
hive> INSERT INTO TRANSACTIONS_HBase
> SELECT * FROM TRANSACTIONS COUNT;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a diff
1.X releases.
Query ID = acadgild_20180424185030_790fe6fd-1a96-4407-8a5e-88518f401e0e
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1524549565701_0021, Tracking URL = http://localhost:8088/proxy/application_1524549565701_0021/
Kill Command = /home/acadgild/install/hadoop/hadoop-2.6.5/bin/hadoop job -kill job_1524549565701_0021
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2018-04-24 18:50:40,648 Stage-3 map = 0%, reduce = 0%
2018-04-24 18:50:49,259 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 4.39 sec
MapReduce Total cumulative CPU time: 4 seconds 390 msec
Ended Job = job_1524549565701_0021
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 4.39 sec HDFS Read: 11269 HDFS Write: 0 SUCCESS
Total MapReduce CPU Time Spent: 4 seconds 390 msec
OK
Time taken: 19.777 seconds

```

scan 'TRANSACTIONS'

Above command will show the content of the table.

Below screen shot shows the content of the TRANSACTIONS table by using scan command

```
acadgild@localhost:~  
hbase(main):025:0> scan 'TRANSACTIONS'  
ROW                                COLUMN+CELL  
101                                column=stats:count_txn, timestamp=1524576048122, value=2  
101                                column=stats:username, timestamp=1524576048122, value=Amitabh  
102                                column=stats:count_txn, timestamp=1524576048122, value=1  
102                                column=stats:username, timestamp=1524576048122, value=Sharukh  
104                                column=stats:count_txn, timestamp=1524576048122, value=1  
104                                column=stats:username, timestamp=1524576048122, value=Anubhav  
105                                column=stats:count_txn, timestamp=1524576048122, value=1  
105                                column=stats:username, timestamp=1524576048122, value=Pawan  
106                                column=stats:count_txn, timestamp=1524576048122, value=1  
106                                column=stats:username, timestamp=1524576048122, value=Aamir  
107                                column=stats:count_txn, timestamp=1524576048122, value=1  
107                                column=stats:username, timestamp=1524576048122, value=Salman  
108                                column=stats:count_txn, timestamp=1524576048122, value=1  
108                                column=stats:username, timestamp=1524576048122, value=Randhir  
7 row(s) in 0.0380 seconds
```

6. Now from the Hbase level, write the Hbase java API code to access and scan the TRANSACTIONS table data from java level.

To solve above problem two-java program coded in the eclipse platform to scan and access the Transaction table.

Program to access the hbase table

```
import java.io.IOException;
```

```
import org.apache.hadoop.conf.Configuration;
```

```
import org.apache.hadoop.hbase.HBaseConfiguration;
```

```
import org.apache.hadoop.hbase.client.Get;
```

```
import org.apache.hadoop.hbase.client.HTable;
```

```
import org.apache.hadoop.hbase.client.Result;
```

```
import org.apache.hadoop.hbase.util.Bytes;
```

```
public class accessHbaseTable{
```

```
    public static void main(String[] args) throws IOException, Exception{
```

```
        // Instantiating Configuration class
```

```
        Configuration config = HBaseConfiguration.create();
```



```

// Instantiating HTable class

@SuppressWarnings({ "resource", "deprecation" })

    HTable table = new HTable(config, "TRANSACTIONS");

// Instantiating Get class

    Get g = new Get(Bytes.toBytes("101"));

// Reading the data

    Result result = table.get(g);


// Reading values from Result class object

    byte [] name = result.getValue(Bytes.toBytes("stats"),Bytes.toBytes("username"));

    byte [] txn = result.getValue(Bytes.toBytes("stats"),Bytes.toBytes("count_txn"));


// Printing the values

    String user = Bytes.toString(name);

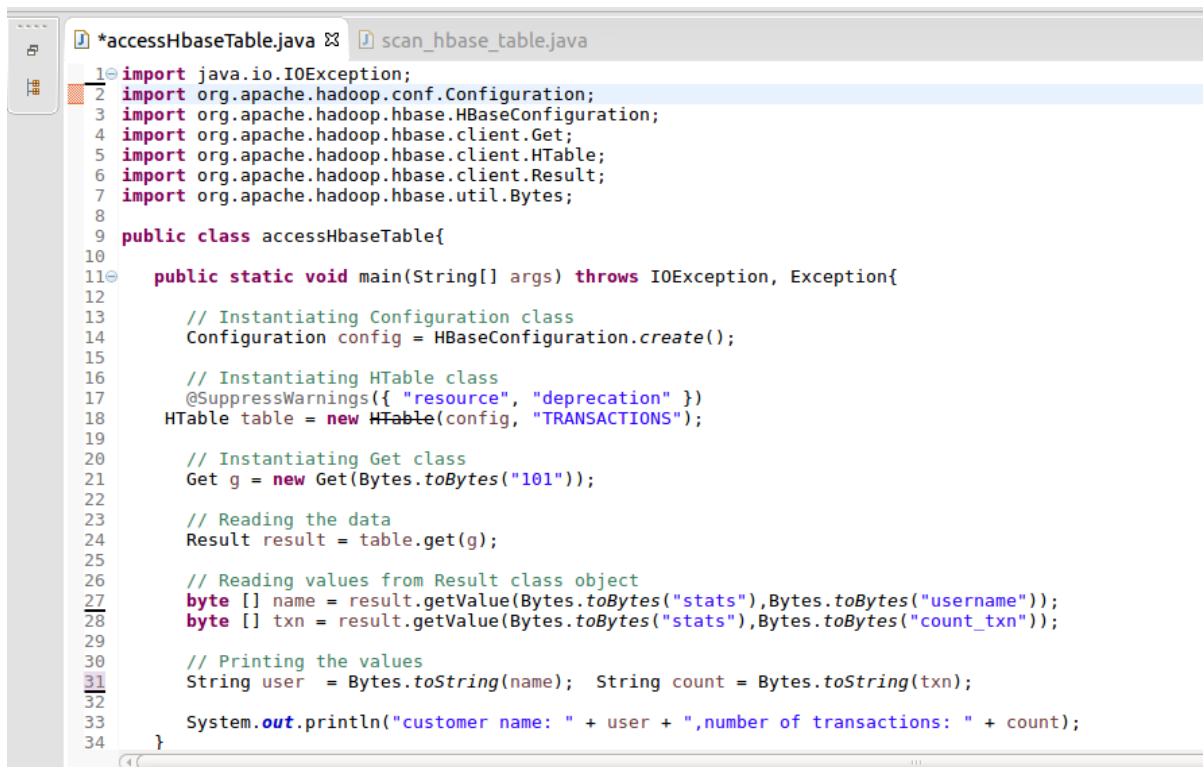
    String count = Bytes.toString(txn);

    System.out.println("customer name: " + user + ",number of transactions: " + count);

}

}

```



```

1 import java.io.IOException;
2 import org.apache.hadoop.conf.Configuration;
3 import org.apache.hadoop.hbase.HBaseConfiguration;
4 import org.apache.hadoop.hbase.client.Get;
5 import org.apache.hadoop.hbase.client.HTable;
6 import org.apache.hadoop.hbase.client.Result;
7 import org.apache.hadoop.hbase.util.Bytes;
8
9 public class accessHbaseTable{
10
11     public static void main(String[] args) throws IOException, Exception{
12
13         // Instantiating Configuration class
14         Configuration config = HBaseConfiguration.create();
15
16         // Instantiating HTable class
17         @SuppressWarnings({ "resource", "deprecation" })
18         HTable table = new HTable(config, "TRANSACTIONS");
19
20         // Instantiating Get class
21         Get g = new Get(Bytes.toBytes("101"));
22
23         // Reading the data
24         Result result = table.get(g);
25
26         // Reading values from Result class object
27         byte [] name = result.getValue(Bytes.toBytes("stats"),Bytes.toBytes("username"));
28         byte [] txn = result.getValue(Bytes.toBytes("stats"),Bytes.toBytes("count_txn"));
29
30         // Printing the values
31         String user = Bytes.toString(name); String count = Bytes.toString(txn);
32
33         System.out.println("customer name: " + user + ",number of transactions: " + count);
34     }
35 }

```

Output: Access table program shows the value of row key 101

```
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details
2018-05-16 14:11:49,488 WARN [main] util.NativeCodeLoader (NativeCodeLoader.java:45) - Unable to load native code library: /lib/x86_64-linux-gnu/libc.so.6
2018-05-16 14:11:49,704 INFO [main] zookeeper.RecoverableZooKeeper (RecoverableZooKeeper.java:100) - ZooKeeper state: Syncing
customer name: Amitabh,number of transactions: 2
```

Program to scan the hbase TRANSACTIONS table:

```
import java.io.IOException;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.hbase.HBaseConfiguration;
import org.apache.hadoop.hbase.util.Bytes;
import org.apache.hadoop.hbase.client.HTable;
import org.apache.hadoop.hbase.client.Result;
import org.apache.hadoop.hbase.client.ResultScanner;
import org.apache.hadoop.hbase.client.Scan;

public class scan_hbase_table{

    public static void main(String args[]) throws IOException{

        // Instantiating Configuration class
        Configuration config = HBaseConfiguration.create();

        // Instantiating HTable class
        @SuppressWarnings({ "deprecation", "resource" })

        HTable table = new HTable(config, "TRANSACTIONS");

        // Instantiating the Scan class
```

```

Scan scan = new Scan();

// scanning the required columns
scan.addColumn(Bytes.toBytes("stats"), Bytes.toBytes("count_txn"));
scan.addColumn(Bytes.toBytes("stats"), Bytes.toBytes("username"));

// Getting the scan result
ResultScanner scanner = table.getScanner(scan);

// Reading values from scan result
for (Result result = scanner.next(); result != null; result = scanner.next())

{
    //assign row values in variable Row
    String Row = Bytes.toString(result.getRow());

    //assign column username values in name
    String name = Bytes.toString(result.getValue("stats".getBytes(), "username".getBytes()));

    //assign column count_txn values in count
    String count = Bytes.toString(result.getValue("stats".getBytes(), "count_txn".getBytes()));

    System.out.println( Row + "," + name + "," + count );

    //closing the scanner
    scanner.close();

}
}

```

```

1 import java.io.IOException;
2
3 import org.apache.hadoop.conf.Configuration;
4 import org.apache.hadoop.hbase.HBaseConfiguration;
5 import org.apache.hadoop.hbase.util.Bytes;
6 import org.apache.hadoop.hbase.client.HTable;
7 import org.apache.hadoop.hbase.client.Result;
8 import org.apache.hadoop.hbase.client.ResultScanner;
9 import org.apache.hadoop.hbase.client.Scan;
10
11
12 public class scan_hbase_table{
13
14     public static void main(String args[]) throws IOException{
15
16         // Instantiating Configuration class
17         Configuration config = HBaseConfiguration.create();
18
19         // Instantiating HTable class
20         @SuppressWarnings({ "deprecation", "resource" })
21
22         HTable table = new HTable(config, "TRANSACTIONS");
23
24         // Instantiating the Scan class
25         Scan scan = new Scan();
26
27         // Scanning the required columns
28         scan.addColumn(Bytes.toBytes("stats"), Bytes.toBytes("count_txn"));
29         scan.addColumn(Bytes.toBytes("stats"), Bytes.toBytes("username"));
30
31         // Getting the scan result
32         ResultScanner scanner = table.getScanner(scan);
33
34         // Getting the scan result
35         ResultScanner scanner = table.getScanner(scan);
36
37         // Reading values from scan result
38         for (Result result = scanner.next(); result != null; result = scanner.next())
39         {
40             //assign row values in variable Row
41             String Row = Bytes.toString(result.getRow());
42
43             //assign column username values in name
44             String name = Bytes.toString(result.getValue(Bytes.toBytes("stats"), Bytes.toBytes("username")));
45
46             //assign column count_txn values in count
47             String count = Bytes.toString(result.getValue(Bytes.toBytes("stats"), Bytes.toBytes("count_txn")));
48
49             System.out.println( Row + "," + name + "," + count );
50
51             //closing the scanner
52             scanner.close();
53         }
54     }
55 }

```

Output: scan program shows the content of the TRANSACTIONS table

```

Console
<terminated> scan_hbase_table [Java Application] /usr/lib/jvm/java-8-openjdk-amd64/bin/java (16-May
SLF4J: Failed to load class "org.slf4j.impl.StaticLoggerBinder".
SLF4J: Defaulting to no-operation (NOP) logger implementation
SLF4J: See http://www.slf4j.org/codes.html#StaticLoggerBinder for further details.
2018-05-16 14:14:35,811 WARN [main] util.NativeCodeLoader (NativeCodeLoader.java:<clin
2018-05-16 14:14:36,057 INFO [main] zookeeper.RecoverableZooKeeper (RecoverableZooKeep
101,Amitabh,2
102,Sharukh,1
104,Anubhav,1
105,Pawan,1
106,Aamir,1
107,Salman,1
108,Ranbir,1

```